

MACHINE LEARNING

Submitted by,

- Name:Jagannath v v
- Reg.No.-21122024
- class:1st MSc DataScience

Lab Overview

Objectives

Questions:

1. Demonstrate the Logistic Regression for different penalties/regularisation methods - none, l1, l2 (you may use 'saga' solver as the parameter)
2. What happens when the Maximum Iterations are kept as 1, 2, 5, 10, 20, 50, 100, 500 and 1000? Is there any change in the accuracy.
3. Get the attributes: classes, coef and intercept_ and print the same in the above case.

Problem Definition

Apply Logistic Regression for Breast Cancer Dataset.

In [1]:

```
import numpy as np
import pandas as pd
import seaborn as sns
```

Performing Logistic Regression

Loading dataset

In [2]:

```
df=pd.read_csv(r'C:\Users\jagan\OneDrive\Desktop\data.csv')
df.head()
```

Out[2]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	co
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	co
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	

5 rows × 33 columns

--	--	--

In [3]:

`df.tail()`

Out[3]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	co
564	926424	M	21.56	22.39	142.00	1479.0	0.11100	
565	926682	M	20.13	28.25	131.20	1261.0	0.09780	
566	926954	M	16.60	28.08	108.30	858.1	0.08455	
567	927241	M	20.60	29.33	140.10	1265.0	0.11780	
568	92751	B	7.76	24.54	47.92	181.0	0.05263	

5 rows × 33 columns

--	--	--

In [4]:

`df.info()`

class 'pandas.core.frame.DataFrame'			
RangeIndex: 569 entries, 0 to 568			
Data columns (total 33 columns):			
#	Column	Non-Null Count	Dtype
0	id	569 non-null	int64
1	diagnosis	569 non-null	object
2	radius_mean	569 non-null	float64
3	texture_mean	569 non-null	float64
4	perimeter_mean	569 non-null	float64
5	area_mean	569 non-null	float64
6	smoothness_mean	569 non-null	float64
7	compactness_mean	569 non-null	float64
8	concavity_mean	569 non-null	float64
9	concave points_mean	569 non-null	float64
10	symmetry_mean	569 non-null	float64
11	fractal_dimension_mean	569 non-null	float64
12	radius_se	569 non-null	float64
13	texture_se	569 non-null	float64
14	perimeter_se	569 non-null	float64
15	area_se	569 non-null	float64
16	smoothness_se	569 non-null	float64
17	compactness_se	569 non-null	float64
18	concavity_se	569 non-null	float64
19	concave points_se	569 non-null	float64
20	symmetry_se	569 non-null	float64
21	fractal_dimension_se	569 non-null	float64
22	radius_worst	569 non-null	float64

```
23  texture_worst      569 non-null   float64
24  perimeter_worst    569 non-null   float64
25  area_worst         569 non-null   float64
26  smoothness_worst   569 non-null   float64
27  compactness_worst  569 non-null   float64
28  concavity_worst   569 non-null   float64
29  concave points_worst  569 non-null   float64
30  symmetry_worst    569 non-null   float64
31  fractal_dimension_worst  569 non-null   float64
32  Unnamed: 32        0 non-null    float64
dtypes: float64(31), int64(1), object(1)
memory usage: 146.8+ KB
```

In [5]: `df.shape`

Out[5]: (569, 33)

In [6]: *### to check the missing values*
`df.isna().sum()`

```
id                      0
diagnosis                0
radius_mean               0
texture_mean              0
perimeter_mean            0
area_mean                 0
smoothness_mean           0
compactness_mean          0
concavity_mean            0
concave points_mean       0
symmetry_mean             0
fractal_dimension_mean    0
radius_se                 0
texture_se                0
perimeter_se              0
area_se                   0
smoothness_se             0
compactness_se            0
concavity_se              0
concave points_se          0
symmetry_se               0
fractal_dimension_se       0
radius_worst               0
texture_worst              0
perimeter_worst            0
area_worst                 0
smoothness_worst           0
compactness_worst          0
concavity_worst            0
concave points_worst        0
symmetry_worst             0
fractal_dimension_worst     0
Unnamed: 32                  569
dtype: int64
```

In [7]: `df.dropna(axis=1, inplace=True)`

#

label encoding

```
In [8]: df['diagnosis'].value_counts
```

```
Out[8]: <bound method IndexOpsMixin.value_counts of 0      M
1      M
2      M
3      M
4      M
..
564    M
565    M
566    M
567    M
568    B
Name: diagnosis, Length: 569, dtype: object>
```

```
In [9]: df.dtypes
```

```
Out[9]: id                  int64
diagnosis          object
radius_mean        float64
texture_mean       float64
perimeter_mean    float64
area_mean          float64
smoothness_mean   float64
compactness_mean  float64
concavity_mean    float64
concave_points_mean float64
symmetry_mean     float64
fractal_dimension_mean float64
radius_se          float64
texture_se         float64
perimeter_se       float64
area_se            float64
smoothness_se     float64
compactness_se    float64
concavity_se      float64
concave_points_se float64
symmetry_se       float64
fractal_dimension_se float64
radius_worst       float64
texture_worst      float64
perimeter_worst   float64
area_worst         float64
smoothness_worst  float64
compactness_worst float64
concavity_worst   float64
concave_points_worst float64
symmetry_worst    float64
fractal_dimension_worst float64
dtype: object
```

Normalizing the labels

```
In [10]: from sklearn.preprocessing import LabelEncoder
```

```
In [11]: labelencoder = LabelEncoder()
df.iloc[:,1] = labelencoder.fit_transform(df.iloc[:,1].values)
```

Splitting the data and feature scaling

```
In [12]: X = df.iloc[:,2:].values
y = df.iloc[:,1].values
```

```
In [13]: from sklearn.model_selection import train_test_split
```

```
In [14]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.40)
```

```
In [15]: #importing standard scalar
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

```
In [16]: X_train
```

```
Out[16]: array([[-0.77791603,  2.27198102, -0.82261281, ..., -1.67376413,
       -2.14202434, -1.39317017],
      [-0.44199969,  0.55408636, -0.4396182 , ...,  0.16612432,
       0.12354534,  0.57114016],
      [ 0.54837433,  0.80598373,  0.65008764, ...,  1.43190651,
       0.61017553,  2.47867512],
      ...,
      [-0.56072874, -0.99814606, -0.54682306, ..., -0.94605389,
       0.36847179, -0.40691781],
      [ 0.45570776,  0.02759817,  0.48192316, ...,  0.63792975,
       -0.09398802, -0.15581855],
      [-1.05012409, -1.61994677, -1.02945513, ..., -0.65136636,
       0.0671478 , -0.3296565 ]])
```

```
In [17]: X_test
```

```
Out[17]: array([[-0.84062568, -0.76506367, -0.86960634, ..., -0.8673476 ,
       -0.0942035 ,  0.10625965],
      [-0.12192206,  1.92987885, -0.12837143, ..., -0.26462778,
       -1.07249206, -0.12301821],
      [-1.04119413, -1.02195567, -1.06177835, ..., -1.34789862,
       -0.99110366, -0.92808438],
      ...,
      [-0.53698733, -0.02831679, -0.51029312, ..., -0.45087005,
       -1.28247413, -0.71073726],
      [-0.72362741, -0.20280948, -0.69600557, ..., -0.12667055,
```

```
0.79455782, -0.10019417],
[-1.32811844, -0.2343151 , -1.36578156, ..., -1.87106309,
-1.62756095, -1.0370173 ]])
```

In [18]:

y_train

Out[18]:

```
array([0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0,
1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0,
0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0,
0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0,
0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0,
0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1,
0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
```

In [19]:

y_test

Out[19]:

```
array([0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1,
0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1,
0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1,
0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0,
0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1,
0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0,
```

Logistic Regression Model

In [20]:

```
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression()
classifier.fit(X_train, y_train)
```

Out[20]:

LogisticRegression()

In [21]:

prediction

In [22]:

predictions = classifier.predict(X_test)

In [23]:

predictions

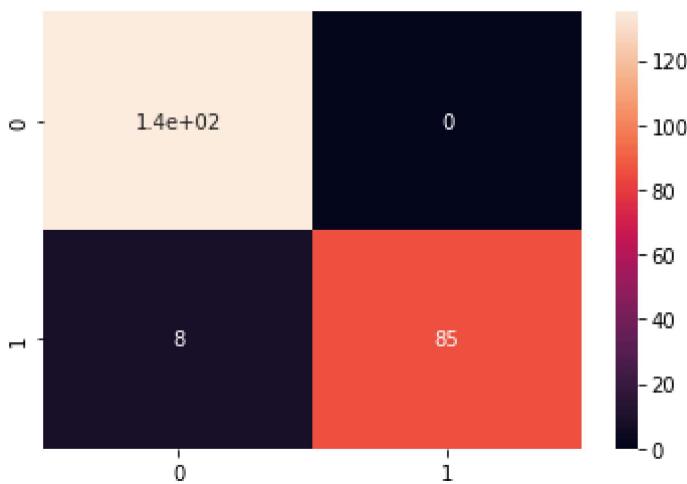
Out[23]:

```
array([0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1,
0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1,
```

```
In [24]: from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, predictions)  
print(cm)  
sns.heatmap(cm, annot=True)
```

```
[[135    0]
 [ 8   85]]
```

Out[24]: <AxesSubplot:>



In [25]: # Accuracy

```
In [26]: from sklearn.metrics import accuracy_score  
print(accuracy_score(y_test,predictions))
```

0.9649122807017544

1. Demonstrate the Logistic Regression for different penalties/regularisation methods - none, l1, l2 (you may use 'saga' solver as the parameter)

```
In [28]: x1_classifier=LogisticRegression(penalty='l2', solver='saga')  
x1_classifier.fit(X_train,y_train)  
predictions2 = x1_classifier.predict(X_test)  
print(accuracy_score(y_test,predictions2))
```

0.9649122807017544

```
C:\Users\jagan\anaconda3\lib\site-packages\sklearn\linear_model\_sag.py:328: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
    warnings.warn("The max_iter was reached which means "
```

In [29]:

```
x1_classifier=LogisticRegression(penalty='l1', solver='saga')
x1_classifier.fit(X_train,y_train)
predictions2 = x1_classifier.predict(X_test)
print(accuracy_score(y_test,predictions2))
```

0.9692982456140351

C:\Users\jagan\anaconda3\lib\site-packages\sklearn\linear_model_sag.py:328: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
 warnings.warn("The max_iter was reached which means "

In [30]:

```
x1_classifier=LogisticRegression(penalty='none', solver='saga')
x1_classifier.fit(X_train,y_train)
predictions2 = x1_classifier.predict(X_test)
print(accuracy_score(y_test,predictions2))
```

0.9649122807017544

C:\Users\jagan\anaconda3\lib\site-packages\sklearn\linear_model_sag.py:328: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
 warnings.warn("The max_iter was reached which means "

2. What happens when the Maximum Iterations are kept as 1, 2, 5, 10, 20, 50, 100, 500 and 1000? Is there any change in the accuracy.

In [38]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

li_iter =[1,2,5,10,20,50,100,500,1000]
x_df = pd.DataFrame()
x_df['Max_iteration'] = [1,2,5,10,20,50,100,500,1000]
li2 = []
li3 = []
li4 = []
li5 = []

j = 1
for i in li_iter:
    classifier = LogisticRegression(penalty = 'l1', max_iter=i, solver = 'saga')
    classifier.fit(X_train, y_train)
    predictions = classifier.predict(X_test)
    li2.append(accuracy_score(y_test,predictions))
    li3.append(classifier.classes_)
    li4.append(classifier.coef_)
    li5.append(classifier.intercept_)
    j += 1

x_df['Accuracy'] = li2
x_df['Classes'] = li3
x_df["Coefficient"] = li4
x_df["Intercept"] = li5

accu_df
```

C:\Users\jagan\anaconda3\lib\site-packages\sklearn\linear_model_sag.py:328: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
 warnings.warn("The max_iter was reached which means "

C:\Users\jagan\anaconda3\lib\site-packages\sklearn\linear_model_sag.py:328: Convergence

Out[38]:	Max_iteration	Accuracy	Classes	Coefficient	Intercept
0	1	0.978070	[0, 1]	[[0.22640988246097374, 0.1930797054922324, 0.2...]	[-0.19238080933748627]
1	2	0.978070	[0, 1]	[[0.29891250199731695, 0.23481326756328422, 0....]	[-0.24766165273982962]
2	5	0.973684	[0, 1]	[[0.35551620312369603, 0.26804835240620944, 0....]	[-0.3868290550403477]
3	10	0.969298	[0, 1]	[[0.3635969578840166, 0.32086308535804864, 0.3...]	[-0.5476265045624693]
4	20	0.973684	[0, 1]	[[0.3862842192267156, 0.37555310772941053, 0.3...]	[-0.6310504100822649]
5	50	0.973684	[0, 1]	[[0.34434875800227394, 0.37231648316902954, 0....]	[-0.7769959307366661]
6	100	0.973684	[0, 1]	[[0.25211005706414935, 0.29944271048523974, 0....]	[-0.8026969258485244]
7	500	0.964912	[0, 1]	[[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.379981447641...]	[-0.7681990786587571]
8	1000	0.964912	[0, 1]	[[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.195209878073...]	[-0.80047927117984491]

3. Get the attributes: `classes`, `coef` and `intercept_` and print the same in the above case.

```
In [33]: x1 classifier classes
```

```
Out[33]: array([0, 1])
```

```
In [37]: x1_classification_coeff
```

```
Out[37]: array([[ 0.32397834,  0.47490267,  0.31838556,  0.39995586,  0.13572746,
   -0.20244219,  0.73389684,  0.83025999,  0.08282472, -0.38353385,
   1.14493002, -0.13397226,  0.79948955,  0.803119 , -0.04129248,
  -0.7928175 , -0.17153889,  0.17295983, -0.45751415, -0.55126836,
   0.83769465,  1.03575343,  0.72762918,  0.80981679,  0.84653306,
   0.08998799,  0.76970489,  0.86912781,  0.59765169,  0.29508445]])
```

```
In [36]: x1_classifier.intercept_
```

```
Out[36]: array([-0.89014637])
```

Conclusion

```
In [ ]: here i learned more about logistic regression
```

References

1. <https://www.quora.com/What-is-regularization-in-machine-learning>
2. <https://towardsdatascience.com/regularization-in-machine-learning-76441ddcf99a>
3. <https://towardsdatascience.com/regularization-an-important-concept-in-machine-learning-5891628907ea>
4. <https://www.geeksforgeeks.org/regularization-in-machine-learning/>
5. <https://analyticsindiamag.com/regularization-in-machine-learning-a-detailed-guide/>
6. <https://towardsdatascience.com/l1-and-l2-regularization-methods-ce25e7fc831c>

```
In [ ]:
```